

Hidden Markov Models

Part 2: Algorithms

CSE 4309 – Machine Learning
Vassilis Athitsos
Computer Science and Engineering Department
University of Texas at Arlington

Hidden Markov Model

- An HMM consists of:
 - A set of states s_1, \dots, s_K .
 - An **initial state probability** function $\pi_k = p(z_1 = s_k)$.
 - A **state transition matrix** A , of size $K \times K$, where

$$A_{k,j} = p(z_n = s_j \mid z_{n-1} = s_k)$$

- **Observation probability functions**, also called **emission probabilities**, defined as:

$$\varphi_k(x) = p(x_n = x \mid z_n = s_k)$$

The Basic HMM Problems

- We will next see algorithms for the four problems that we usually want to solve when using HMMs.
- Problem 1: learn an HMM using training data. This involves:
 - Learning the initial state probabilities π_k .
 - Learning the state transition matrix A .
 - Learning the observation probability functions $P_k(x)$.
- Problems 2, 3, 4 assume we have already learned an HMM.
- Problem 2: given a sequence of observations $X = x_1, \dots, x_N$, compute $p(X \mid \text{HMM})$: the probability of X given the model.
- Problem 3: given an observation sequence $X = x_1, \dots, x_N$, find the hidden state sequence $Z = z_1, \dots, z_N$ maximizing $p(Z \mid X, \text{HMM})$
- Problem 4: given an observation sequence $X = x_1, \dots, x_N$, compute, for any n, k , the probability $p(z_n = s_k \mid X, \text{HMM})$

The Basic HMM Problems

- Problem 1: learn an HMM using training data.
- Problem 2: given a sequence of observations $X = x_1, \dots, x_N$, compute $p(X \mid \text{HMM})$: the probability of X given the model.
- Problem 3: given an observation sequence $X = x_1, \dots, x_N$, find the hidden state sequence $Z = z_1, \dots, z_N$ maximizing $p(Z \mid X, \text{HMM})$
- Problem 4: given an observation sequence $X = x_1, \dots, x_N$, compute, for any n, k , the probability $p(z_n = s_k \mid X, \text{HMM})$
- We will first look at algorithms for problems 2, 3, 4.
- Last, we will look at the standard algorithm for learning an HMM using training data.

Probability of Observations

- Problem 2: given an HMM model θ , and a sequence of observations $X = x_1, \dots, x_N$, compute $p(X \mid \theta)$: the probability of X given the model.
- Inputs:
 - $\theta = (\pi, A, \varphi_k)$: a trained HMM, with probability functions π, A, φ_k specified.
 - A sequence of observations $X = x_1, \dots, x_N$.
- Output: $p(X \mid \theta)$.

Probability of Observations

- Problem 2: given a sequence of observations $X = x_1, \dots, x_N$, compute $p(X \mid \theta)$.
- Why do we care about this problem?
 - What would be an example application?

Probability of Observations

- Problem 2: given a sequence of observations $X = x_1, \dots, x_N$, compute $p(X | \theta)$.
- Why do we care about this problem?
- We need to compute $p(X | \theta)$ to classify X .
 - Suppose we have multiple models $\theta_1, \dots, \theta_c$.
 - For example, we can have one model for digit 2, one model for digit 3, one model for digit 4...
 - A Bayesian classifier classifies X by finding the θ_c that maximizes $p(\theta_c | X)$.
 - Using Bayes rule, we must maximize $\frac{p(X | \theta_c) p(\theta_c)}{p(X)}$.
 - Therefore, we need to compute $p(X | \theta_c)$ for each c .

The Sum Rule

- The sum rule, which we saw earlier in the course, states that:

$$p(X) = \sum_{y \in \mathbb{Y}} p(X, Y = y)$$

- We can use the sum rule, to compute $p(X \mid \theta)$ as follows:

$$p(X \mid \theta) = \sum_Z p(X, Z \mid \theta)$$

- According to this formula, we can compute $p(X \mid \theta)$ by summing $p(X, Z \mid \theta)$ over **all possible state sequences Z** .

The Sum Rule

$$p(X | \theta) = \sum_Z p(X, Z | \theta)$$

- According to this formula, we can compute $p(X | \theta)$ by summing $p(X, Z | \theta)$ over **all possible state sequences** Z .
- An HMM with parameters θ specifies a joint distribution function $p(X, Z | \theta)$ as:

$$p(X, Z | \theta) = p(z_1 | \theta) \prod_{n=2}^N p(z_n | z_{n-1}, \theta) \prod_{n=1}^N p(x_n | z_n, \theta)$$

- Therefore:

$$p(X | \theta) = \sum_Z \left[p(z_1 | \theta) \prod_{n=2}^N p(z_n | z_{n-1}, \theta) \prod_{n=1}^N p(x_n | z_n, \theta) \right]$$

The Sum Rule

$$p(X \mid \theta) = \sum_Z \left[p(z_1 \mid \theta) \prod_{n=2}^N p(z_n \mid z_{n-1}, \theta) \prod_{n=1}^N p(x_n \mid z_n, \theta) \right]$$

- According to this formula, we can compute $p(X \mid \theta)$ by summing $p(X, Z \mid \theta)$ over **all possible state sequences Z** .
- What would be the time complexity of doing this computation directly?
- It would be linear to the number of all possible state sequences Z , which is typically exponential to N (the length of X).
- Luckily, there is a polynomial time algorithm for computing $p(X \mid \theta)$, that uses dynamic programming.

Dynamic Programming for $p(X | \theta)$

$$p(X | \theta) = \sum_Z \left[p(z_1 | \theta) \prod_{n=2}^N p(z_n | z_{n-1}, \theta) \prod_{n=1}^N p(x_n | z_n, \theta) \right]$$

- To compute $p(X | \theta)$, we use a dynamic programming algorithm that is called the **forward algorithm**.
- We define a 2D array of problems, of size $N \times K$.
 - N : length of observation sequence X .
 - K : number of states in the HMM.
- Problem (n, k) :
 - Compute $\alpha[n, k] = p(x_1, \dots, x_n, z_n = s_k | \theta)$

The Forward Algorithm - Initialization

$$p(X | \theta) = \sum_Z \left[p(z_1 | \theta) \prod_{n=2}^N p(z_n | z_{n-1}, \theta) \prod_{n=1}^N p(x_n | z_n, \theta) \right]$$

- Problem (n, k) :
 - Compute $\alpha[n, k] = p(x_1, \dots, x_n, z_n = s_k | \theta)$
- Problem $(1, k)$:
 - ???

The Forward Algorithm - Initialization

$$p(X | \theta) = \sum_Z \left[p(z_1 | \theta) \prod_{n=2}^N p(z_n | z_{n-1}, \theta) \prod_{n=1}^N p(x_n | z_n, \theta) \right]$$

- Problem (n, k) :
 - Compute $\alpha[n, k] = p(x_1, \dots, x_n, z_n = s_k | \theta)$
- Problem $(1, k)$:
 - Compute $\alpha[1, k] = p(x_1, z_1 = s_k | \theta)$
- Solution:
 - ???

The Forward Algorithm - Initialization

$$p(X | \theta) = \sum_Z \left[p(z_1 | \theta) \prod_{n=2}^N p(z_n | z_{n-1}, \theta) \prod_{n=1}^N p(x_n | z_n, \theta) \right]$$

- Problem (n, k) :
 - Compute $\alpha[n, k] = p(x_1, \dots, x_n, z_n = s_k | \theta)$
- Problem $(1, k)$:
 - Compute $\alpha[1, k] = p(x_1, z_1 = s_k | \theta)$
- Solution:
 - $\alpha[1, k] = p(x_1, z_1 = k | \theta) = \pi_k \varphi_k(x_1)$

The Forward Algorithm - Main Loop

- Problem (n, k) :
 - Compute $\alpha[n, k] = p(x_1, \dots, x_n, z_n = s_k \mid \theta)$
- Solution:

The Forward Algorithm - Main Loop

- Problem (n, k) :
 - Compute $\alpha[n, k] = p(x_1, \dots, x_n, z_n = s_k | \theta)$
- Solution:

$$p(x_1, \dots, x_n, z_n = s_k | \theta) = \sum_{j=1}^K p(x_1, \dots, x_n, z_{n-1} = s_j, z_n = s_k | \theta)$$

$$= \varphi_k(x_n) \sum_{j=1}^K p(x_1, \dots, x_{n-1}, z_{n-1} = s_j, z_n = s_k | \theta)$$

$$= \varphi_k(x_n) \sum_{j=1}^K \{p(x_1, \dots, x_{n-1}, z_{n-1} = s_j | \theta) A_{j,k}\}$$

The Forward Algorithm - Main Loop

- Problem (n, k) :
 - Compute $\alpha[n, k] = p(x_1, \dots, x_n, z_n = s_k | \theta)$
- Solution:

$$p(x_1, \dots, x_n, z_n = k | \theta)$$

$$= \varphi_k(x_n) \sum_{j=1}^K \{p(x_1, \dots, x_{n-1}, z_{n-1} = s_j | \theta) A_{j,k}\}$$

$$= \varphi_k(x_n) \sum_{j=1}^K \{\alpha[n-1, j] A_{j,k}\}$$

The Forward Algorithm - Main Loop

- Problem (n, k) :
 - Compute $\alpha[n, k] = p(x_1, \dots, x_n, z_n = s_k | \theta)$
- Solution:

$$p(x_1, \dots, x_n, z_n = k | \theta) = \varphi_k(x_n) \sum_{j=1}^K \{\alpha[n-1, j] A_{j,k}\}$$

- Thus, $\alpha[n, k] = \varphi_k(x_n) \sum_{j=1}^K \{\alpha[n-1, j] A_{j,k}\}$.
- $\alpha[n, k]$ is easy to compute once all $\alpha[n-1, j]$ values have been computed.

The Forward Algorithm

- Our aim was to compute $p(X | \theta)$.
- Using the dynamic programming algorithm we just described, we can compute $\alpha[N, k] = p(X, z_N = s_k | \theta)$.
- How can we use those $\alpha[N, k]$ values to compute $p(X | \theta)$?
- We use the sum rule:

$$p(X | \theta) = \sum_{k=1}^K \{p(X, z_N = s_k | \theta)\} \Rightarrow$$

$$p(X | \theta) = \sum_{k=1}^K \alpha[n, k]$$

Problem 3: Find the Best Z

- Problem 3:
 - Inputs: a trained HMM θ , and an observation sequence $X = x_1, \dots, x_N$
 - Output: the state sequence $Z = z_1, \dots, z_N$ maximizing $p(Z | \theta, X)$.
- Why do we care?
- Some times, we want to use HMMs for classification.
 - In those cases, we do not really care about maximizing $p(Z | \theta, X)$, we just want to find the HMM parameters θ that maximize $p(\theta | X)$.
- Some times, we want to use HMMs to figure out the most likely values of the hidden states given the observations.
 - In the tree ring example, our goal was to figure out the average temperature for each year.
 - Those average temperatures were the hidden states.
 - Solving problem 3 provides the most likely sequence of average temperatures.

Problem 3: Find the Best Z

- Problem 3:
 - Input: an observation sequence $X = x_1, \dots, x_N$
 - Output: the state sequence $Z = z_1, \dots, z_N$ maximizing $p(Z \mid \theta, X)$.
- We want to maximize $p(Z \mid \theta, X)$.
- Using the definition of conditional probabilities:

$$p(Z \mid \theta, X) = \frac{p(X, Z \mid \theta)}{p(X \mid \theta)}$$

- Since X is known, $p(X \mid \theta)$ will be the same over all possible state sequences Z .
- Therefore, to find the Z that maximizes $p(Z \mid \theta, X)$, it suffices to find the Z that maximizes $p(Z, X \mid \theta)$.

Problem 3: Find the Best Z

- Problem 3:
 - Input: an observation sequence $X = x_1, \dots, x_N$
 - Output: the state sequence $Z = z_1, \dots, z_N$ maximizing $p(Z \mid \theta, X)$, which is the same as maximizing $p(Z, X \mid \theta)$.

- Solution: (again) dynamic programming.

- Problem (n, k) : Compute $G[n, k]$ and $H[n, k]$, where:

$$G[n, k] = \operatorname{argmax}_{z_1, \dots, z_{n-1}} p(x_1, \dots, x_n, z_1, \dots, z_{n-1}, z_n = s_k \mid \theta)$$

- Note that:
 - We optimize over all possible values z_1, \dots, z_{n-1} .
 - However, **z_n is constrained to be equal to s_k** .
 - Values x_1, \dots, x_n are known, and thus they are fixed.

Problem 3: Find the Best Z

- Problem 3:
 - Input: an observation sequence $X = x_1, \dots, x_N$
 - Output: the state sequence $Z = z_1, \dots, z_N$ maximizing $p(Z \mid \theta, X)$, which is the same as maximizing $p(Z, X \mid \theta)$.
- Solution: (again) dynamic programming.
- Problem (n, k) : Compute $G[n, k]$ and $H[n, k]$, where:

$$H[n, k] = p(x_1, \dots, x_n, G[n, k] \mid \theta)$$

- $H[n, k]$ is the joint probability of:
 - the first n observations x_1, \dots, x_n
 - the sequence we stored in $G[n, k]$

Problem 3: Find the Best Z

- Problem 3:
 - Input: an observation sequence $X = x_1, \dots, x_N$
 - Output: the state sequence $Z = z_1, \dots, z_N$ maximizing $p(Z \mid \theta, X)$, which is the same as maximizing $p(Z, X \mid \theta)$.
- Solution: (again) dynamic programming.
- Problem (n, k) : Compute $G[n, k]$ and $H[n, k]$, where:

$$H[n, k] = p(x_1, \dots, x_n, G[n, k] \mid \theta)$$

- $H[n, k]$ is the joint probability of:
 - the first n observations x_1, \dots, x_n
 - the sequence we stored in $G[n, k]$
- $G[n, k]$ stores a **sequence** of state values.
- $H[n, k]$ stores a **number** (a probability).

The Viterbi Algorithm

- Problem (n, k) : Compute $G[n, k]$ and $H[n, k]$, where:

$$G[n, k] = \operatorname{argmax}_{z_1, \dots, z_{n-1}} p(x_1, \dots, x_n, z_1, \dots, z_{n-1}, z_n = s_k | \theta)$$

$$H[n, k] = p(x_1, \dots, x_n, G[n, k] | \theta)$$

- We use a dynamic programming algorithm to compute $G[n, k]$ and $H[n, k]$ for all n, k such that $1 \leq n \leq N, 1 \leq k \leq K$.
- This dynamic programming algorithm is called the **Viterbi algorithm**.

The Viterbi Algorithm - Initialization

- Problem $(1, k)$: Compute $G[1, k]$ and $H[1, k]$, where:

$$G[1, k] = \operatorname{argmax}_{z_1, \dots, z_{n-1}} p(x_1, z_1 = s_k | \theta)$$

$$H[1, k] = p(x_1, z_1 = s_k | \theta)$$

- What is $G[1, k]$?
 - It is the **empty sequence**.
 - We must maximize over the previous states z_1, \dots, z_{n-1} , but $n = 1$, so there are no previous states.
- What is $H[1, k]$?
 - $H[1, k] = p(x_1, z_1 = s_k | \theta) = \pi_k \varphi_k(x_1)$.

The Viterbi Algorithm – Main Loop

- Problem (n, k) : Compute $G[n, k]$ and $H[n, k]$, where:

$$G[n, k] = \operatorname{argmax}_{z_1, \dots, z_{n-1}} p(x_1, \dots, x_n, z_1, \dots, z_{n-1}, z_n = s_k \mid \theta)$$

$$H[n, k] = p(x_1, \dots, x_n, G[n, k] \mid \theta)$$

- How do we find $G[n, k]$?
- $G[n, k] = G^*[n, k] \otimes s_j$ for some $G^*[n, k]$ and some j .
 - The last element of $G[n, k]$ is some s_j .
 - Therefore, $G[n, k]$ is formed by appending some s_j to some sequence $G^*[n, k]$

The Viterbi Algorithm – Main Loop

- Problem (n, k) : Compute $G[n, k]$ and $H[n, k]$, where:

$$G[n, k] = \operatorname{argmax}_{z_1, \dots, z_{n-1}} p(x_1, \dots, x_n, z_1, \dots, z_{n-1}, z_n = s_k \mid \theta)$$

$$H[n, k] = p(x_1, \dots, x_n, G[n, k] \mid \theta)$$

- $G[n, k] = G^*[n, k] \otimes s_j$ for some $G^*[n, k]$ and some j .
- We can prove that $G^*[n, k] = G[n-1, j]$, for that j .
 - The proof is very similar to the proof that DTW finds the optimal alignment.
 - If $G[n-1, j]$ is better than $G^*[n, k]$, then $G[n-1, j] \otimes s_j$ is better than $G[n, k]$, which is a contradiction.

The Viterbi Algorithm – Main Loop

- Problem (n, k) : Compute $G[n, k]$ and $H[n, k]$, where:

$$G[n, k] = \operatorname{argmax}_{z_1, \dots, z_{n-1}} p(x_1, \dots, x_n, z_1, \dots, z_{n-1}, z_n = s_k \mid \theta)$$

$$H[n, k] = p(x_1, \dots, x_n, G[n, k] \mid \theta)$$

- $G[n, k] = G[n - 1, j] \otimes s_j$ for some j .
- Let's define j^* as: $j^* = \operatorname{argmax}_{j=1, \dots, K} \{H[n - 1, j] A_{j,k} \varphi_k(x_n)\}$
- Then:
 - $G[n, k] = G[n - 1, j^*] \otimes s_{j^*}$.
 - $H[n, k] = \{H[n - 1, j^*] A_{j^*,k} \varphi_k(x_n)\}$

The Viterbi Algorithm – Output

- Our goal is to find the Z maximizing $p(Z \mid \theta, X)$, which is the same as finding the Z maximizing $p(Z, X \mid \theta)$.
- The Viterbi algorithm computes $G[n, k]$ and $H[n, k]$, where:

$$G[n, k] = \operatorname{argmax}_{z_1, \dots, z_{n-1}} p(x_1, \dots, x_n, z_1, \dots, z_{n-1}, z_n = s_k \mid \theta)$$

$$H[n, k] = p(x_1, \dots, x_n, G[n, k] \mid \theta)$$

- Let's define k^* as: $k^* = \operatorname{argmax}_{k=1, \dots, K} \{H[N, k]\}$
- Then, **the Z maximizing $p(Z, X \mid \theta)$ is $G[N, k^*] \otimes s_{k^*}$**

State Probabilities at Specific Times

- Problem 4:
 - Inputs: a trained HMM θ , and an observation sequence $X = x_1, \dots, x_N$
 - Output: an array γ , of size $N \times K$, where $\gamma[n, k] = p(z_n = s_k \mid X, \theta)$.
- In words: given an observation sequence X , we want to compute, for any moment in time n and any state s_k , the probability that the hidden state at that moment is s_k .

State Probabilities at Specific Times

- Problem 4:
 - Inputs: a trained HMM θ , and an observation sequence $X = x_1, \dots, x_N$
 - Output: an array γ , of size $N \times K$, where $\gamma[n, k] = p(z_n = s_k \mid X, \theta)$.
- We have seen, in our solution for problem 2, that the **forward algorithm** computes an array α , of size $N \times K$, where $\alpha[n, k] = p(x_1, \dots, x_n, z_n = s_k \mid \theta)$.
- We can also define another array β , of size $N \times K$, where $\beta[n, k] = p(x_{n+1}, \dots, x_N \mid \theta, z_n = s_k)$.
- In words, $\beta[n, k]$ is the probability of all observations **after** moment n , given that the state at moment n is s_k .

State Probabilities at Specific Times

- Problem 4:
 - Inputs: a trained HMM θ , and an observation sequence $X = x_1, \dots, x_N$
 - Output: an array γ , of size $N \times K$, where $\gamma[n, k] = p(z_n = s_k \mid X, \theta)$.
- We have seen, in our solution for problem 2, that the **forward algorithm** computes an array α , of size $N \times K$, where $\alpha[n, k] = p(x_1, \dots, x_n, z_n = s_k \mid \theta)$.
- We can also define another array β , of size $N \times K$, where $\beta[n, k] = p(x_{n+1}, \dots, x_N \mid \theta, z_n = s_k)$.
- Then, $\alpha[n, k] * \beta[n, k] =$
$$p(x_1, \dots, x_n, z_n = s_k \mid \theta) * p(x_{n+1}, \dots, x_N \mid \theta, z_n = s_k) =$$
$$p(X, z_n = s_k \mid \theta)$$
- Therefore: $\gamma[n, k] = p(z_n = s_k \mid X, \theta) = \frac{\alpha[n, k] * \beta[n, k]}{P(X \mid \theta)}$

State Probabilities at Specific Times

- The **forward algorithm** computes an array α , of size $N \times K$, where $\alpha[n, k] = p(x_1, \dots, x_n, z_n = s_k \mid \theta)$.
- We define another array β , of size $N \times K$, where $\beta[n, k] = p(x_{n+1}, \dots, x_N \mid \theta, z_n = s_k)$.
- Then, $\gamma[n, k] = p(z_n = s_k \mid X, \theta) = \frac{\alpha[n, k] * \beta[n, k]}{P(X \mid \theta)}$
- In the above equation:
 - $\alpha[n, k]$ and $P(X \mid \theta)$ are computed by the forward algorithm.
 - We need to compute $\beta[n, k]$.
- We compute $\beta[n, k]$ in a way very similar to the forward algorithm, but going backwards this time.
 - The resulting combination of the forward and backward algorithms is called (not surprisingly) the **forward-backward algorithm**.

The Backward Algorithm

- We want to compute the values of an array β , of size $N \times K$, where $\beta[n, k] = p(x_{n+1}, \dots, x_N \mid \theta, z_n = s_k)$.
- Again, we will use dynamic programming.
 - Problem (n, k) is to compute value $\beta[n, k]$.
- However, this time we start from the end of the observations.
- First, compute $\beta[N, k] = p(\{\} \mid \theta, z_N = s_k)$.
- In this case, the observation sequence x_{n+1}, \dots, x_N is empty, because $n = N$.
- What is the probability of an empty sequence?

Backward Algorithm - Initialization

- We want to compute the values of an array β , of size $N \times K$, where $\beta[n, k] = p(x_{n+1}, \dots, x_N \mid \theta, z_n = s_k)$.
- Again, we will use dynamic programming.
 - Problem (n, k) is to compute value $\beta[n, k]$.
- However, this time we start from the end of the observations.
- First, compute $\beta[N, k] = p(\{\} \mid \theta, z_N = s_k)$.
- In this case, the observation sequence x_{n+1}, \dots, x_N is empty, because $n = N$.
- What is the probability of an empty sequence?
 - $\beta[N, k] = 1$

Backward Algorithm - Initialization

- Next: compute $\beta[N - 1, k]$.

$$\beta[N - 1, k] = p(x_N | \theta, z_{N-1} = s_k) =$$

$$\sum_{j=1}^K p(x_N, z_N = s_j | \theta, z_{N-1} = s_k) =$$

$$\sum_{j=1}^K \{p(x_N | z_N = s_j)p(z_N = s_j | \theta, z_{N-1} = s_k)\}$$

Backward Algorithm - Initialization

- Next: compute $\beta[N - 1, k]$.

$$\beta[N - 1, k] = p(x_N | \theta, z_{N-1} = s_k)$$

$$= \sum_{j=1}^K \{p(x_N | z_N = s_j) p(z_N = s_j | \theta, z_{N-1} = s_k)\}$$

$$= \sum_{j=1}^K \{\varphi_j(x_N) A_{k,j}\}$$

Backward Algorithm – Main Loop

- Next: compute $\beta[n, k]$, for $n < N - 1$.

$$\beta[n, k] = p(x_{n+1}, \dots, x_N | \theta, z_n = s_k) =$$

$$\sum_{j=1}^K p(x_{n+1}, \dots, x_N, z_{n+1} = s_j | \theta, z_n = s_k) =$$

$$\sum_{j=1}^K \{p(x_{n+1} | z_{n+1} = s_j) p(x_{n+2}, \dots, x_N, z_{n+1} = s_j | \theta, z_{N-1} = s_k)\}$$

Backward Algorithm – Main Loop

- Next: compute $\beta[n, k]$, for $n < N - 1$.

$$\beta[n, k] = p(x_{n+1}, \dots, x_N | \theta, z_n = s_k) =$$

$$\sum_{j=1}^K \{p(x_{n+1} | z_{n+1} = s_j) p(x_{n+2}, \dots, x_N, z_{n+1} = s_j | \theta, z_{N-1} = s_k)\}$$

$$\sum_{j=1}^K \{\varphi_j(x_{n+1}) p(x_{n+2}, \dots, x_N | \theta, z_{n+1} = s_j) A_{k,j}\}$$

We will take a closer look at the last step...

Backward Algorithm – Main Loop

$$\sum_{j=1}^K \{p(x_{n+1} \mid z_{n+1} = s_j) p(x_{n+2}, \dots, x_N, z_{n+1} = s_j \mid \theta, z_n = s_k)\} =$$

$$\sum_{j=1}^K \{\varphi_j(x_{n+1}) p(x_{n+2}, \dots, x_N \mid \theta, z_{n+1} = s_j) A_{k,j}\}$$

- $p(x_{n+1} \mid z_{n+1} = s_j) = \varphi_j(x_{n+1})$, by definition of φ_j .

Backward Algorithm – Main Loop

$$\sum_{j=1}^K \{p(x_{n+1} \mid z_{n+1} = s_j) p(x_{n+2}, \dots, x_N, z_{n+1} = s_j \mid \theta, z_n = s_k)\} =$$

$$\sum_{j=1}^K \{\varphi_j(x_{n+1}) p(x_{n+2}, \dots, x_N \mid \theta, z_{n+1} = s_j) A_{k,j}\}$$

$$p(x_{n+2}, \dots, x_N, z_{n+1} = s_j \mid \theta, z_n = s_k) = \\ p(x_{n+2}, \dots, x_N \mid \theta, z_n = s_k, z_{n+1} = s_j) p(z_{n+1} = s_j \mid z_n = s_k)$$

by application of chain rule: $P(A, B) = P(A|B)P(B)$,

which can also be written as $P(A, B \mid C) = P(A|B, C)P(B \mid C)$

Backward Algorithm – Main Loop

$$\sum_{j=1}^K \{p(x_{n+1} \mid z_{n+1} = s_j) p(x_{n+2}, \dots, x_N, z_{n+1} = s_j \mid \theta, z_n = s_k)\} =$$

$$\sum_{j=1}^K \{\varphi_j(x_{n+1}) p(x_{n+2}, \dots, x_N \mid \theta, z_{n+1} = s_j) A_{k,j}\}$$

$$\begin{aligned} p(x_{n+2}, \dots, x_N, z_{n+1} = s_j \mid \theta, z_n = s_k) &= \\ p(x_{n+2}, \dots, x_N \mid \theta, z_n = s_k, z_{n+1} = s_j) p(z_{n+1} = s_j \mid z_n = s_k) &= \\ p(x_{n+2}, \dots, x_N \mid \theta, z_n = s_k, z_{n+1} = s_j) A_{k,j} \end{aligned}$$

(by definition of $A_{k,j}$)

Backward Algorithm – Main Loop

$$\sum_{j=1}^K \{p(x_{n+1} \mid z_{n+1} = s_j) p(x_{n+2}, \dots, x_N, z_{n+1} = s_j \mid \theta, z_n = s_k)\} =$$

$$\sum_{j=1}^K \{\varphi_j(x_{n+1}) p(x_{n+2}, \dots, x_N \mid \theta, z_{n+1} = s_j) A_{k,j}\}$$

$$\begin{aligned} p(x_{n+2}, \dots, x_N, z_{n+1} = s_j \mid \theta, z_n = s_k) &= \\ p(x_{n+2}, \dots, x_N \mid \theta, z_n = s_k, z_{n+1} = s_j) p(z_{n+1} = s_j \mid z_n = s_k) &= \\ p(x_{n+2}, \dots, x_N \mid \theta, z_n = s_k, z_{n+1} = s_j) A_{k,j} &= \\ p(x_{n+2}, \dots, x_N \mid \theta, z_{n+1} = s_j) A_{k,j} \end{aligned}$$

$(x_{n+2}, \dots, x_N$ are conditionally independent of z_n given $z_{n+1})$

Backward Algorithm – Main Loop

- In the previous slides, we showed that

$$\beta[n, k] = p(x_{n+1}, \dots, x_N | \theta, z_n = s_k) = \sum_{j=1}^K \{ \varphi_j(x_{n+1}) p(x_{n+2}, \dots, x_N | \theta, z_{n+1} = s_j) A_{k,j} \}$$

- Note that $p(x_{n+2}, \dots, x_N | \theta, z_{n+1} = s_j)$ is $\beta[n + 1, j]$.
- Consequently, we obtain the recurrence:

$$\beta[n, k] = \sum_{j=1}^K \{ \varphi_j(x_{n+1}) \beta[n + 1, j] A_{k,j} \}$$

Backward Algorithm – Main Loop

$$\beta[n, k] = \sum_{j=1}^K \{ \varphi_j(x_{n+1}) \beta[n+1, j] A_{k,j} \}$$

- Thus, we can easily compute all values $\beta[n, k]$ using this recurrence.
- The key thing is to proceed in decreasing order of n , since values $\beta[n, *]$ depend on values $\beta[n+1, *]$.

The Forward-Backward Algorithm

- Inputs:
 - A trained HMM θ .
 - An observation sequence $X = x_1, \dots, x_N$
- Output:
 - An array γ , of size $N \times K$, where $\gamma[n, k] = p(z_n = s_k \mid X, \theta)$.
- Algorithm:
 - Use the forward algorithm to compute $\alpha[n, k]$.
 - Use the backward algorithm to compute $\beta[n, k]$.
 - For $n = 1$ to N :
 - For $k = 1$ to K :

$$\gamma[n, k] = p(z_n = s_k \mid X, \theta) = \frac{\alpha[n, k] * \beta[n, k]}{P(X \mid \theta)}$$

Problem 1: Training an HMM

- Goal: Estimate parameters $\theta = (\pi_k, A_{k,j}, \varphi_k)$.
- The training data can consist of multiple observation sequences $X_1, X_2, X_3, \dots, X_M$.
- We denote the length of training sequence X_i as N_m .
- We denote the elements of each observation sequence X_m as:

$$X_m = (x_{m,1}, x_{m,2}, \dots, x_{m,N_m})$$

- **Before** we start training, we need to decide on:
 - The number of states.
 - The transitions that we will allow.

Problem 1: Training an HMM

- While we are given $X_1, X_2, X_3, \dots, X_M$, we are **not** given the corresponding hidden state sequences $Z_1, Z_2, Z_3, \dots, Z_M$.
- We denote the elements of each hidden state sequence Z_m as:
$$Z_m = (z_{m,1}, z_{m,2}, \dots, z_{m,N_m})$$
- The training algorithm is called **Baum-Welch algorithm**.
- It is an Expectation-Maximization algorithm, similar to the algorithm we saw for learning Gaussian mixtures.

Expectation-Maximization

- When we wanted to learn a mixture of Gaussians, we had the following problem:
 - If we knew the probability of each object belonging to each Gaussian, we could estimate the parameters of each Gaussian.
 - If we knew the parameters of each Gaussian, we could estimate the probability of each object belonging to each Gaussian.
 - However, we know neither of these pieces of information.
- The EM algorithm resolved this problem using:
 - An initialization of Gaussian parameters to some random or non-random values.
 - A main loop where:
 - The current values of the Gaussian parameters are used to estimate new weights of membership of every training object to every Gaussian.
 - The current estimated membership weights are used to estimate new parameters (mean and covariance matrix) for each Gaussian.

Expectation-Maximization

- When we want to learn an HMM model using observation sequences as training data, we have the following problem:
 - If we knew, for each observation sequence, the probabilities of the hidden state values, we could estimate the parameters θ .
 - If we knew the parameters θ , we could estimate the probabilities of the hidden state values.
 - However, we know neither of these pieces of information.
- The Baum-Welch algorithm resolves this problem using EM:
 - At initialization, parameters θ are given (mostly) random values.
 - In the main loop, these two steps are performed repeatedly:
 - The current values of parameters θ are used to estimate new probabilities for the hidden state values.
 - The current probabilities for the hidden state values are used to estimate new values for parameters θ .

Baum-Welch: Initialization

- As we said before, before we start training, we need to decide on:
 - The number of states.
 - The transitions that we will allow.
- When we start training, we initialize $\theta = (\pi_k, A_{k,j}, \varphi_k)$ to random values, with these exceptions:
 - For any s_k that we do **not** want to allow to ever be an initial state, we set the corresponding π_k to 0.
 - The rest of the training will keep these π_k values always equal to 0.
 - For any s_k, s_j such that we do **not** want to allow transitions from s_k to s_j , we set the corresponding $A_{k,j}$ to 0.
 - The rest of the training will keep these $A_{k,j}$ values always equal to 0.

Baum-Welch: Initialization

- When we start training, we initialize $\theta = (\pi_k, A_{k,j}, \varphi_k)$ to random values, with these exceptions:
 - For any s_k that we do **not** want to allow to ever be an initial state, we set the corresponding π_k to 0.
 - The rest of the training will keep these π_k values always equal to 0.
 - For any s_k, s_j such that we do **not** want to allow transitions from s_k to s_j , we set the corresponding $A_{k,j}$ to 0.
 - The rest of the training will keep these $A_{k,j}$ values always equal to 0.
- These initial choices constrain the topology of the resulting HMM model.
 - For example, they can force the model to be fully connected, to be a forward model, or to be some other variation.

Baum-Welch: Expectation Step

- Before we start the expectation step, we have some current values for the parameters of $\theta = (\pi_k, A_{k,j}, \varphi_k)$.
- The goal of the expectation step is to use those values to estimate two arrays:
 - $\gamma[m, n, k] = p(x_{m,1}, \dots, x_{m,n}, z_{m,n} = s_k \mid \theta, X_m)$
 - This is the same $\gamma[n, k]$ that we computed earlier, using the forward-backward algorithm.
 - Here we just need to make the array three-dimensional, because we have multiple observation sequences X_m .
 - We can compute $\gamma[m, *, *]$ by running the forward-backward algorithm with θ and X_m as inputs.
 - $\xi[m, n, j, k] = p(z_{m,n-1} = s_j, z_{m,n} = s_k \mid \theta, X_m)$

Baum-Welch: Expectation Step

- To facilitate our computations, we will also extend arrays α and β to three dimensions, in the same way that we extended array γ .
 - $\alpha[m, n, k] = p(x_{m,1}, \dots, x_{m,n}, z_{m,n} = s_k \mid \theta)$
 - $\beta[m, n, k] = p(x_{m,n+1}, \dots, x_{m,N_m} \mid \theta, z_{m,n} = s_k)$
- Where needed, values $\alpha[m, *, *]$, $\beta[m, *, *]$, $\gamma[m, *, *]$ can be computed by running the forward-backward algorithm with inputs θ and X_m .

Computing the ξ Values

- Using Bayes rule:

$$\begin{aligned}\xi[m, n, j, k] &= p(z_{m,n-1} = s_j, z_{m,n} = s_k \mid \theta, X_m) \\ &= \frac{p(X_m \mid \theta, z_{m,n-1} = s_j, z_{m,n} = s_k) p(z_{m,n-1} = s_j, z_{m,n} = s_k \mid \theta)}{p(X_m \mid \theta)}\end{aligned}$$

- $p(X_m \mid \theta)$ can be computed using the forward algorithm.
- We will see how to simplify and compute:
 - $p(X_m \mid \theta, z_{m,n-1} = s_j, z_{m,n} = s_k)$
 - $p(z_{m,n-1} = s_j, z_{m,n} = s_k \mid \theta)$

Computing the ξ Values

$$\begin{aligned}\xi[m, n, j, k] &= p(z_{m,n-1} = s_j, z_{m,n} = s_k \mid \theta, X_m) \\ &= \frac{p(X_m \mid \theta, z_{m,n-1} = s_j, z_{m,n} = s_k) p(z_{m,n-1} = s_j, z_{m,n} = s_k \mid \theta)}{p(X_m \mid \theta)}\end{aligned}$$

$$\begin{aligned}p(X_m \mid \theta, z_{m,n-1} = s_j, z_{m,n} = s_k) &= \\ p(x_{m,1}, \dots, x_{m,n-1} \mid \theta, z_{m,n-1} = s_j) &p(x_{m,n}, \dots, x_{m,N_m} \mid \theta, z_{m,n} = s_k)\end{aligned}$$

- Why?
 - Earlier observations $x_{m,1}, \dots, x_{m,n-1}$ are conditionally independent of state $z_{m,n}$ given state $z_{m,n-1}$.
 - Later observations $x_{m,n}, \dots, x_{m,N_m}$ are conditionally independent of state $z_{m,n-1}$ given state $z_{m,n}$.

Computing the ξ Values

$$\begin{aligned}\xi[m, n, j, k] &= p(z_{m,n-1} = s_j, z_{m,n} = s_k \mid \theta, X_m) \\ &= \frac{p(X_m \mid \theta, z_{m,n-1} = s_j, z_{m,n} = s_k) p(z_{m,n-1} = s_j, z_{m,n} = s_k \mid \theta)}{p(X_m \mid \theta)}\end{aligned}$$

$$\begin{aligned}p(X_m \mid \theta, z_{m,n-1} = s_j, z_{m,n} = s_k) &= \\ p(x_{m,1}, \dots, x_{m,n-1} \mid \theta, z_{m,n-1} = s_j) & p(x_{m,n}, \dots, x_{m,N_m} \mid \theta, z_{m,n} = s_k)\end{aligned}$$

$$\begin{aligned}p(x_{m,n}, \dots, x_{m,N_m} \mid \theta, z_{m,n} = s_k) &= \\ p(x_{m,n} \mid \theta, z_{m,n} = s_k) & p(x_{m,n+1}, \dots, x_{m,N_m} \mid \theta, z_{m,n} = s_k) = \\ \varphi_k(x_{m,n}) \beta[m, n, k]\end{aligned}$$

Computing the ξ Values

$$\begin{aligned}\xi[m, n, j, k] &= p(z_{m,n-1} = s_j, z_{m,n} = s_k \mid \theta, X_m) \\ &= \frac{p(X_m \mid \theta, z_{m,n-1} = s_j, z_{m,n} = s_k) p(z_{m,n-1} = s_j, z_{m,n} = s_k \mid \theta)}{p(X_m \mid \theta)}\end{aligned}$$

$$\begin{aligned}p(X_m \mid \theta, z_{m,n-1} = s_j, z_{m,n} = s_k) &= \\ p(x_{m,1}, \dots, x_{m,n-1} \mid \theta, z_{m,n-1} = s_j) \varphi_k(x_{m,n}) \beta[m, n, k] &= \\ \frac{p(x_{m,1}, \dots, x_{m,n-1}, z_{m,n-1} = s_j \mid \theta)}{p(z_{m,n-1} = s_j \mid \theta)} \varphi_k(x_{m,n}) \beta[m, n, k] &= \\ \frac{\alpha[m, n-1, j]}{\gamma[m, n-1, j]} \varphi_k(x_{m,n}) \beta[m, n, k]\end{aligned}$$

Computing the ξ Values

$$\xi[m, n, j, k] =$$

$$\frac{\frac{\alpha[m, n-1, j]}{\gamma[m, n-1, j]} \varphi_k(x_{m,n}) \beta[m, n, k] p(z_{m,n-1} = s_j, z_{m,n} = s_k | \theta)}{p(X_m | \theta)}$$

- We can further process $p(z_{m,n-1} = s_j, z_{m,n} = s_k | \theta)$:

$$\begin{aligned} p(z_{m,n-1} = s_j, z_{m,n} = s_k | \theta) &= \\ p(z_{m,n} = s_k | \theta, z_{m,n-1} = s_j) p(z_{m,n-1} = s_j | \theta) &= \\ A_{j,k} \gamma[m, n-1, j] \end{aligned}$$

Computing the ξ Values

- So, we get:

$$\xi[m, n, j, k] =$$

$$\frac{\frac{\alpha[m, n - 1, j]}{\gamma[m, n - 1, j]} \varphi_k(x_{m,n}) \beta[m, n, k] A_{j,k} \gamma[m, n - 1, j]}{p(X_m | \theta)} =$$

$$\frac{\alpha[m, n - 1, j] \varphi_k(x_{m,n}) \beta[m, n, k] A_{j,k}}{p(X_m | \theta)}$$

Computing the ξ Values

- Based on the previous slides, the final formula for ξ is:

$$\xi[m, n, j, k] = \frac{\alpha[m, n - 1, j] \varphi_k(x_{m,n}) \beta[m, n, k] A_{j,k}}{p(X_m | \theta)}$$

- This formula can be computed using the parameters θ and the algorithms we have covered:
 - $\alpha[m, n - 1, j]$ is computed with the forward algorithm.
 - $\beta[m, n, k]$ is computed by the backward algorithm.
 - $p(X_m | \theta)$ is computed with the forward algorithm.
 - φ_k and $A_{j,k}$ are specified by θ .

Baum-Welch: Summary of E Step

- Inputs:
 - Training observation sequences $X_1, X_2, X_3, \dots, X_M$.
 - Current values of parameters $\theta = (\pi_k, A_{k,j}, \varphi_k)$.
- Algorithm:
 - For $m = 1$ to M :
 - Run the forward-backward algorithm to compute:
 - $\alpha[m, n, k]$ for n and k such that $1 \leq n \leq N_m, 1 \leq k \leq K$.
 - $\beta[m, n, k]$ for n and k such that $1 \leq n \leq N_m, 1 \leq k \leq K$.
 - $\gamma[m, n, k]$ for n and k such that $1 \leq n \leq N_m, 1 \leq k \leq K$.
 - $$p(X_m | \theta)$$
 - For n, j, k such that $1 \leq n \leq N_m, 1 \leq j, k \leq K$:
$$\xi[m, n, j, k] = \frac{\alpha[m, n-1, j] \varphi_k(x_{m,n}) \beta[m, n, k] A_{j,k}}{p(X_m | \theta)}$$

Baum-Welch: Maximization Step

- Before we start the maximization step, we have some current values for:
 - $\gamma[m, n, k] = p(z_{m,n} = s_k \mid \theta, X_m)$
 - $\xi[m, n, j, k] = p(z_{m,n-1} = s_j, z_{m,n} = s_k \mid \theta, X_m)$
- The goal of the maximization step is to use those values to estimate new values for $\theta = (\pi_k, A_{k,j}, \varphi_k)$.
- In other words, we want to estimate new values for:
 - Initial state probabilities π_k .
 - State transition probabilities $A_{k,j}$.
 - Observation probabilities φ_k .

Baum-Welch: Maximization Step

- Initial state probabilities π_k are computed as:

$$\pi_k = \frac{\sum_{m=1}^M \gamma[m, 1, k]}{\sum_{m=1}^M \sum_{j=1}^K \gamma[m, 1, j]}$$

- In words, we compute for each k the ratio of these two quantities:
 - The sum of probabilities, over all observation sequences X_m , that state s_k was the initial state for X_m .
 - The sum of probabilities over all observation sequences X_m **and all states** s_j that states s_j was the initial state for X_m .

Baum-Welch: Maximization Step

- State transition probabilities $A_{k,j}$ are computed as:

$$A_{k,j} = \frac{\sum_{m=1}^M \sum_{n=2}^{N_m} \xi[m, n, k, j]}{\sum_{m=1}^M \sum_{n=2}^{N_m} \sum_{i=1}^K \xi[m, n, k, i]}$$

- In words, we compute, for each k, j the ratio of these two quantities:
 - The sum of probabilities, over all state transitions of all observation sequences X_m , that state s_k was followed by s_j .
 - The sum of probabilities, over all state transitions of all observation sequences X_m **and all states** s_i , that state s_k was followed by state s_i .

Baum-Welch: Maximization Step

- Computing the observation probabilities φ_k depends on how we model those probabilities.
- For example, we could choose:
 - Discrete probabilities.
 - Histograms.
 - Gaussians.
 - Mixtures of Gaussians.
 - ...
- We will show how to compute distributions φ_k for the cases of:
 - Discrete probabilities.
 - Gaussians.

Baum-Welch: Maximization Step

- Computing the observation probabilities φ_k depends on how we model those probabilities.
- In all cases, the key idea is that we treat observation $x_{m,n}$ as partially assigned to distribution φ_k .
 - $\gamma[m, n, k] = p(z_{m,n} = s_k \mid \theta, X_m)$ is the weight of the assignment of $x_{m,n}$ to φ_k .
- In other words:
 - We don't know what hidden state $x_{m,n}$ corresponds to, but we have computed, for each state s_k , the probability $\gamma[m, n, k]$ that $x_{m,n}$ corresponds to s_k .
 - Thus, $x_{m,n}$ influences our estimate of distribution φ_k with weight proportional to $\gamma[m, n, k]$.

Baum-Welch: Maximization Step

- Suppose that the observations are discrete, and come from a finite set $\mathbf{Y} = \{y_1, \dots, y_R\}$.
 - In that case, each $x_{m,n}$ is an element of \mathbf{Y} .
- Define an auxiliary function $\text{Eq}(x, y)$ as:

$$\text{Eq}(x, y) = \begin{cases} 1 & \text{if } x = y \\ 0 & \text{if } x \neq y \end{cases}$$

- Then, $\varphi_k(y_r) = p(x_{m,n} = y_r \mid z_{m,n} = s_k)$, and it can be computed as:

$$\varphi_k(y_r) = \frac{\sum_{m=1}^M \sum_{n=1}^{N_m} \{\gamma[m, n, k] \text{Eq}(x_{m,n}, y_r)\}}{\sum_{m=1}^M \sum_{n=1}^{N_m} \gamma[m, n, k]}$$

Baum-Welch: Maximization Step

- If observations come from a finite set $Y = \{y_1, \dots, y_R\}$:

$$\varphi_k(y_r) = \frac{\sum_{m=1}^M \sum_{n=1}^{N_m} \{\gamma[m, n, k] \text{Eq}(x_{m,n}, y_r)\}}{\sum_{m=1}^M \sum_{n=1}^{N_m} \gamma[m, n, k]}$$

- The above formula can be seen as a weighted average of the times $x_{m,n}$ is equal to y_r , where the weight of each $x_{m,n}$ is the probability that $x_{m,n}$ corresponds to hidden state s_k .

Baum-Welch: Maximization Step

- Suppose that observations are vectors in \mathbb{R}^D , and that we model φ_k as a Gaussian distribution.
- In that case, for each φ_k we need to estimate a mean μ_k and a covariance matrix S_k .

$$\mu_k(y_r) = \frac{\sum_{m=1}^M \sum_{n=1}^{N_m} \{\gamma[m, n, k] x_{m,n}\}}{\sum_{m=1}^M \sum_{n=1}^{N_m} \gamma[m, n, k]}$$

- Thus, μ_k is a weighted average of values $x_{m,n}$, where the weight of each $x_{m,n}$ is the probability that $x_{m,n}$ corresponds to hidden state s_k .

Baum-Welch: Maximization Step

- Suppose that observations are vectors in \mathbb{R}^D , and that we model φ_k as a Gaussian distribution.
- In that case, for each φ_k we need to estimate a mean μ_k and a covariance matrix S_k .

$$\mu_k = \frac{\sum_{m=1}^M \sum_{n=1}^{N_m} \{\gamma[m, n, k] x_{m,n}\}}{\sum_{m=1}^M \sum_{n=1}^{N_m} \gamma[m, n, k]}$$

$$S_k = \frac{\sum_{m=1}^M \sum_{n=1}^{N_m} \left\{ \gamma[m, n, k] (x_{m,n} - \mu_k)(x_{m,n} - \mu_k)^T \right\}}{\sum_{m=1}^M \sum_{n=1}^{N_m} \gamma[m, n, k]}$$

Baum-Welch: Summary of M Step

$$\pi_k = \frac{\sum_{m=1}^M \gamma[m, 1, k]}{\sum_{m=1}^M \sum_{j=1}^K \gamma[m, 1, j]} \quad A_{k,j} = \frac{\sum_{m=1}^M \sum_{n=2}^{N_m} \xi[m, n, k, j]}{\sum_{m=1}^M \sum_{n=2}^{N_m} \sum_{i=1}^K \xi[m, n, k, i]}$$

- For discrete observations:

$$\varphi_k(y_r) = \frac{\sum_{m=1}^M \sum_{n=1}^{N_m} \{\gamma[m, n, k] \text{Eq}(x_{m,n}, y_r)\}}{\sum_{m=1}^M \sum_{n=1}^{N_m} \gamma[m, n, k]}$$

Baum-Welch: Summary of M Step

$$\pi_k = \frac{\sum_{m=1}^M \gamma[m, 1, k]}{\sum_{m=1}^M \sum_{j=1}^K \gamma[m, 1, j]} \quad A_{k,j} = \frac{\sum_{m=1}^M \sum_{n=2}^{N_m} \xi[m, n, k, j]}{\sum_{m=1}^M \sum_{n=2}^{N_m} \sum_{i=1}^K \xi[m, n, k, i]}$$

- For Gaussian observation distributions: φ_k is a Gaussian with mean μ_k and covariance matrix S_k , where:

$$\mu_k = \frac{\sum_{m=1}^M \sum_{n=1}^{N_m} \{\gamma[m, n, k] x_{m,n}\}}{\sum_{m=1}^M \sum_{n=1}^{N_m} \gamma[m, n, k]}$$

$$S_k = \frac{\sum_{m=1}^M \sum_{n=1}^{N_m} \{\gamma[m, n, k] (x_{m,n} - \mu_k)(x_{m,n} - \mu_k)^T\}}{\sum_{m=1}^M \sum_{n=1}^{N_m} \gamma[m, n, k]}$$

Baum-Welch Summary

- Initialization:
 - Initialize $\theta = (\pi_k, A_{k,j}, \varphi_k)$ with random values, but set to 0 values π_k and $A_{k,j}$ to specify the network topology.
- Main loop:
 - E-step:
 - Compute all values $p(X_m | \theta)$, $\alpha[m, n, k]$, $\beta[m, n, k]$, $\gamma[m, n, k]$ using the forward-backward algorithm.
 - Compute all values $\xi[m, n, j, k] = \frac{\alpha[m, n-1, j] \varphi_k(x_{m,n}) \beta[m, n, k] A_{j,k}}{p(X_m | \theta)}$.
 - M-step:
 - Update $\theta = (\pi_k, A_{k,j}, \varphi_k)$, using the values $\gamma[m, n, k]$ and $\xi[m, n, j, k]$ computed at the E-step.

Hidden Markov Models - Recap

- HMMs are widely used to model temporal sequences.
- In HMMs, an observation sequence corresponds to a sequence of hidden states.
- An HMM is defined by parameters $\theta = (\pi_k, A_{k,j}, \varphi_k)$, and defines a joint probability $p(X, Z | \theta)$.
- An HMM can be used as a generative model, to produce synthetic samples from distribution $p(X, Z | \theta)$.
- HMMs can be used for Bayesian classification:
 - One HMM θ_c per class.
 - Find the class c that maximizes $p(\theta_c | X)$.

Hidden Markov Models - Recap

- HMMs can be used for Bayesian classification:
 - One HMM θ_c per class.
 - Find the class c that maximizes $p(\theta_c | X)$, using probabilities $p(X | \theta_c)$ calculated by the forward algorithm.
- HMMs can be used to find the most likely hidden state sequence Z for observation sequence X .
 - This is done using the Viterbi algorithm.
- HMMs can be used to find the most likely hidden state z_n for a specific value x_n of an observation sequence X .
 - This is done using the forward-backward algorithm.
- HMMs are trained using the Baum-Welch algorithm, which is an Expectation-Maximization algorithm.